

# Ticket Whisperer

AI support-ticket triage that uses **the right model for the right job** — and proves the savings in real time.

Track: Multi-model

Solo build


[ticket-whisperer.onrender.com](https://ticket-whisperer.onrender.com)

Smiten Kalathiya · [smiten.kalathiya6@gmail.com](mailto:smiten.kalathiya6@gmail.com)

## Support triage burns money **twice**

- ▶ **Human time.** Tier-1 agents spend minutes per ticket figuring out what it's about, whether it was solved before, and how to respond — before writing a single word.
- ▶ **Model waste.** Teams that automate this point one flagship LLM at everything. But classification is a cheap-model task, retrieval is an embeddings task — only the customer-facing draft needs premium reasoning.
- ▶ **Result:** paying flagship prices for work a model 20× cheaper does just as well.

# One pipeline, three models, one gateway



Every call routed through **Mesh API** · one base URL · one key · per-call cost tracking

## Why **Mesh** makes this possible

### One API, 500+ models

OpenAI + Anthropic + hundreds more behind a single OpenAI-compatible endpoint. Our three-provider pipeline is one base URL and one key — no SDK juggling.

### Live cost transparency

The app pulls live pricing from Mesh's `/models` endpoint and computes the real cost of every step from actual token usage. No hardcoded price tables.

### Instant model swapping

Every step's model is an env var. Try a different classifier or drafter by changing one line in `.env` — no code changes, no new SDKs.

# The cost meter **proves it live**

| STEP         | MODEL VIA MESH         | COST / ANALYSIS |
|--------------|------------------------|-----------------|
| Classify     | gpt-4o-mini            | ~\$0.0001       |
| Retrieve     | text-embedding-3-small | ~\$0.00001      |
| Draft        | claude-sonnet-4.6      | ~\$0.0009       |
| <b>Total</b> | smart routing          | ~\$0.001        |
|              | everything on flagship | ~\$0.01+        |

# ~90%

cheaper per ticket —  
computed from **real token usage**  
and shown on screen after every run

# Built lean

- ▶ **Node.js / Express** — thin orchestrator; dependencies are express + dotenv only.
- ▶ **Zero database.** Knowledge base is a JSON file, embedded at startup via Mesh embeddings, held in memory.
- ▶ **SSE streaming** — each stage streams status, output, and cost to the browser live, driving the pipeline animation and cost meter.
- ▶ **BYO-key UI** — judges paste their own Mesh key on the page; it stays in localStorage, never stored or logged server-side.
- ▶ **Mesh integration:** `lib/mesh.js` (sole AI client) · `lib/pipeline.js` (3-step triage) · `lib/pricing.js` (live /models pricing)

## Try it right now

Open the live demo, paste your Mesh key, pick a sample ticket, hit ⚡ **Analyze** — and watch three providers cooperate on one ticket while the meter counts the cost.

**Live demo** · <https://ticket-whisperer.onrender.com>

**Video walkthrough** · <https://www.loom.com/share/e6031bedf95646e89d8f0cd255271d44>

**Source** · <https://github.com/sk5202/ticket-whisperer>

Ticket Whisperer — **automate the triage, stop overpaying for it.**